

Andrey Asadchev

CONTACT INFORMATION

(515) 337-5362
asadchev@gmail.com

SUMMARY

Computational scientist passionate about C++, high-performance computing, and low-level code optimization. Interests include C++ template expressions and meta-programming, parallel linear algebra, computational chemistry, and machine learning.

US permanent resident (green card holder).

Fluent in English, native Russian speaker, conversational level of French.

EDUCATION

PhD, Physical Chemistry, Dec 2012

Iowa State University, Ames, IA

- Emphasis on computational quantum chemistry.
- Dissertation: *Modernizing the core quantum chemistry algorithms*,
<http://lib.dr.iastate.edu/etd/12915/>

BSc, Computer Science, May, 2005

Valdosta State University, Valdosta, GA

- Magna Cum Laude

TECHNICAL SKILLS

Programming Languages:	C++/C++11, Python, C, FORTRAN 77/90
C++ Concepts:	Object oriented design, metaprogramming, expression templates
C++ Libraries:	Standard Template Library (STL), Boost (including MPL and Fusion)
Machine Learning:	Tensorflow, Scikit, R, Skytree
Scientific Computing:	BLAS, LAPACK, Eigen, Trillions, Mathematica, Matlab/Octave
Parallel Programming:	MPI, CUDA, TBB, OpenMP, ARMCI
Parallel I/O:	HDF5, NetCDF, Lustre, PVFS
Code optimization:	SSE and AVX vectorization with intrinsics/assembly
Performance analysis:	Linux perf, Intel VTune/Parallel Studio, PAPI
Debugging:	gdb, Valgrind, DDT, TotalView
Build tools:	CMake, autotools
Collaboration:	git, Atlassian stack (JIRA, Stash), Agile
HPC Environments:	IBM BlueGene, Cray XT, OpenMPI, MPICH, PBS/Torque
UNIX/Linux:	POSIX API programming, bash shell scripting, Linux administration

WORK AND RESEARCH EXPERIENCE

Infosys Limited, Palo Alto, CA

Senior Product Technical Architect

2017 - Present

Development and integration of machine learning features in Infosys products

- Integrated Tensorflow into Infosys ML offerings.
- Introduced C++/Python interoperability via Boost.Python to allow for rapid development and prototyping.
- Continuous enhancement of core code base and build system.

Skytree, Inc., San Jose, CA*HPC C++ Developer***2013 - 2017**

Participated in the development of the Skytree C++ Machine Learning (ML) product.

- Refactored the C++ codebase according to modern C++ patterns; reduced the number of lines of code, improved maintainability and compilation speed.
- Identified and resolved key computational and memory bottlenecks; significantly improved overall execution time and reduced memory footprint.
- Implemented core dense and sparse linear algebra objects and kernels.
- Designed and implemented several distributed ML algorithms using OpenMP and MPI.
- Worked with data scientists to address feature requests and resolve bugs.

Virginia Tech, Blacksburg, VA*PostDoctoral Researcher***2012 - 2013**

Participated in the development of Massively Parallel Quantum Chemistry (MPQC) package.

- Modernized the MPQC C++ code base by using Boost and Eigen linear algebra library.
- Ported MPQC to CMake; simplified build and deployment across various HPC platforms.
- Developed fully distributed OpenMP/MPI sparse Davidson eigensolver with out-of-core capabilities, scalable to 1000s of cores.
- Conducted pilot research to optimize key computational kernels on SSE/AVX processors.

Iowa State University, Ames, IA*Research Assistant***2008 - 2012**

Developed computational chemistry C++ library (libcchem) with GPU capabilities and FORTRAN bindings.

- Implemented automatically generated vectorized C++ Rys Quadrature algorithm; 30-40% faster than the original FORTRAN code.
- Implemented scalable multi-threaded Hartree-Fock algorithm with constant memory overhead.
- Implemented Rys Quadrature and Hartree-Fock algorithms on GPU using CUDA C++; showed speedups on the order of 10-17 times.
In the process resolved several issues in Boost C++ libraries due to CUDA compiler.
- Designed and implemented distributed in-core and out-of-core data storage suitable for matrix and tensor computations with C++ object oriented interface and integration with BLAS.
- Implemented scalable distributed perturbation theory and Coupled Cluster algorithms with low memory footprint and ability to utilize GPU devices via CuBLAS.
The implementation showed orders of magnitude reduction in memory use, as well as improvement in speed and overall scalability (1000s of cores), over the original FORTRAN code.

Ames Laboratory, Ames, IA*Research Assistant***2006 - 2008**

Participated in the development of The General Atomic and Molecular Electronic Structure System (GAMESS).

- Added ARMCi remote memory implementation to GAMESS.
- Implemented execution script to run GAMESS in various MPI (OpenMPI, MPICH, IMPI) and batch system environments (LSF, PBS, Torque, LoadLeveler).
- Ported GAMESS to Blue Gene and Cray XT.
- Replaced FORTRAN code with BLAS/LAPACK calls, parallelized linear algebra routines; improved overall program performance, allowed GAMESS to scale to 1000s of cores.

TEACHING EXPERIENCE

Virginia Tech, Blacksburg, VA

(SICM)2 Summer School

2013

- Taught GPU/CUDA programming workshop to computational sciences graduate students.

Iowa State University, Ames, IA

Chemistry Instructor

2005 - 2006

- Taught general chemistry lab and recitation sections.

PUBLICATIONS AND PRESENTATIONS

A Fast and Flexible Coupled Cluster Approach

A. Asadchev, M. S. Gordon

Journal of Chemical Theory and Computation 2013 9(8)

New Multithreaded Hybrid CPU/GPU Approach to HartreeFock

A. Asadchev, M. S. Gordon

Journal of Chemical Theory and Computation 2012 8 (11)

Uncontracted Rys Quadrature Implementation of up to G Functions on graphical processing units (GPUs)

A. Asadchev, V. Allada, J. Felder, B. M. Bode, M. S. Gordon, T. L. Windus

Journal of Chemical Theory and Computation 2010 6(3)

Accelerating Quantum Chemistry Research using GPUs

A. Asadchev, J. Felder

GPU Technology Conference, NVIDIA, San Jose, 2009

Uncontracted Rys Quadrature on GPU

A. Asadchev

Path to Petascale, NCSA, 2009

Performance of Electronic Structure Calculations on Blue Gene/L and Cray XT4

A. Asadchev

IEEE/ACM Supercomputing 2008 Poster